

# The Search-First & Action-Ranked UI Architecture

Author: Antigravity PM Team | Date: September 2, 2026 | Target: PrepBench Next-Gen UI

## 1. Executive Thesis: Dissolving the Taxonomy Trap

For four sessions, the product debate centered on whether PrepBench should organize its interface by **Formats** (Exams, Interviews, System Design) or by **Subjects** (Scrum, Databricks, AI). This proposal asserts that **the two-axis debate is a false dilemma** created entirely by trying to build a traditional navigation tree.

**The Real Google Move:** Google's core architectural signature is not hierarchical taxonomies—it is *refusing to make the user navigate*. Gmail killed folders with search + labels. Drive is a search box with a filesystem attached. Power users live in command palettes. By establishing a single global Search / Command Input (Cmd + K) as the front door, the subject-vs-format debate completely dissolves: searching 'cost' or 'scrum events' instantaneously returns exams, design reviews, guide passages, and personal accuracy in one surface.

## 2. Explicit Assumptions & Risks (For Counter-Argument)

To enable rigorous counter-debate, the core assumptions underlying this design proposal are explicitly stated below alongside their vulnerabilities:

| Assumption                       | Stated Rationale                                                                                             | Counter-Argument / Vulnerability                                                                                                             |
|----------------------------------|--------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------|
| 1. Session Intent                | Users open the app with a specific topic or keyword in mind (e.g. 'I want to study Delta Lake').             | <b>Vulnerability:</b> Novice users or passive learners may suffer from 'empty text box paralysis' when they don't know what keyword to type. |
| 2. Command Palette Preference    | Target users are technical engineers who prefer keyboard shortcuts (Cmd+K) over clicking sidebar links.      | <b>Vulnerability:</b> Mouse-driven users may ignore keyboard shortcuts entirely if visual affordances aren't prominent.                      |
| 3. Action-Ranked Home            | Users prefer 3 ranked actionable recommendations (Resume, SM-2 Review, Weak Spot) over passive KPI charts.   | <b>Vulnerability:</b> Users who want a birds-eye summary of overall certification readiness may feel deprived of visual progress gauges.     |
| 4. SM-2 Spaced Repetition Wiring | Surfacing the existing sm2_service.py backend engine on Home will drive daily short reviews.                 | <b>Vulnerability:</b> If a user misses 2 weeks, a 150-item SM-2 backlog could create overwhelming study debt.                                |
| 5. Inline AI vs Drawer AI        | Placing the [Why?] button directly inline on incorrect questions reduces friction compared to a side-drawer. | <b>Vulnerability:</b> Inline AI expansions can clutter the quiz card layout if detailed multi-paragraph explanations are generated.          |

## 3. UI Architecture & Wireframes

### A. Home Screen: Search-First & Action-Ranked (3 Action Cards)

```
■ ■ Search anything or press Cmd + K (e.g. 'scrum events', 'delta lake', 'cost')... ■
■
■ 1. ■■ RESUME SESSION
■ Databricks Design Review #3 ■■ Step 4 of 6 (Last active 18h ago)
■ [ ■■ Continue Design Review ]
■
■ 2. ■ DUE FOR SPACED REPETITION REVIEW (Powered by sm2_service.py)
■ 14 questions queued based on your memory decay curve (Scrum Events & Delta Lake)
■ [ ■ Start 5-Minute Spaced Review ]
■
■ 3. ■ WEAKEST TOPIC DRILL
■ Storage vs Compute Cost Optimization (38% accuracy across 3 attempts)
■ [ ■ Target This Weakness ]
■
```

B. Command Palette Search Overlay (⌘ + K)

```
■ ■ Query: 'cost'
■
■ ■ DESIGN REVIEW : Databricks FinOps & DBU Cost Optimization (Case Study #4)
■ ■ EXAM QUESTIONS: 5 Questions on DBU Pricing & Cluster Autoscaling
■ ■■ ROADMAP TOPIC : Databricks Governance & Cost Management (Phase 3)
■ ■ YOUR ACCURACY : 38% Accuracy on Cost Questions (3 Missed)
```

C. Practice Surface with Inline AI at Failure Point

```
■ Question 3 of 10 • Databricks FinOps
■ What is the primary advantage of Auto-terminating single-node clusters?
■
■ [ ] A. It increases Spark driver memory allocation.
■ [X] B. It prevents idle DBU consumption when developer notebooks are inactive. (YOUR)
■ [✓] C. It eliminates control plane management overhead. (CORRECT)
■
■ ■ INLINE AI EXPLANATION (One-Click Inline Expansion)
■ Auto-termination stops compute charges, but control plane management remains managed
■ by Databricks. Option C accurately reflects architectural boundaries under DBU rules.
```

4. Backend Engine Integration Plan

A critical finding during this audit: **sm2\_service.py** and the **SpacedRepetition** model exist in the backend codebase, but have zero API router endpoints or frontend connections. The mathematical engine is already built—it simply requires wiring.

| Engine Component       | Current Backend State                                                                             | Proposed Surfacing                                                                        |
|------------------------|---------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------|
| SM-2 Spaced Repetition | Fully implemented in app/services/sm2_service.py with ease factor floors (1.3) and interval math. | Wire to /api/v1/spaced-repetition/due and surface directly as Card #2 on the Home screen. |
| Weak Spot Detection    | Calculated inside analytics_repository.py and system_design_service.py.                           | Surface as Card #3 on Home: 1-click targeted drill for the user's lowest accuracy topic.  |
| LLM Gateway            | Abstracted in app/llm/gateway.py with fallback to local env / Gemini.                             | Wire directly to inline [Why?] feedback buttons on quiz failure cards.                    |

## 5. Empirical Telemetry Plan (2-Week Usage Query)

Rather than debating taxonomy theoretically, PrepBench can settle the Subject vs. Format question empirically. Every practice attempt is logged in SQLite (exam\_simulator.db) with timestamps, domains, and session modes. Running the following SQL query after 2 weeks of natural usage will provide definitive evidence:

```
SELECT
  q.domain,
  s.exam_mode,
  COUNT(*) as total_attempts
FROM exam_answers a
JOIN questions q ON a.question_id = q.id
JOIN exam_sessions s ON a.session_id = s.id
GROUP BY q.domain, s.exam_mode;
```

**Decision Metric:** If usage clusters heavily by domain across multiple modes in a single sitting, subject-first scoping is empirically validated. If usage clusters by exam mode across multiple domains, format-first navigation is validated.